

Reinforcement Learning and Optimal Control

ASU, CSE 691, Winter 2019

Dimitri P. Bertsekas
dimitrib@mit.edu

Lecture 3

- 1 Approximation in Value and Policy Space
- 2 General Issues of Approximation in Value Space
- 3 Special Multistep Lookahead Issues
- 4 Problem Approximation Schemes

Recall the Stochastic DP Algorithm

Produces the optimal costs $J_k^*(x_k)$ of the tail subproblems that start at x_k

Start with $J_N^*(x_N) = g_N(x_N)$, and for $k = 0, \dots, N - 1$, let

$$J_k^*(x_k) = \min_{u_k \in U_k(x_k)} E \left\{ g_k(x_k, u_k, w_k) + J_{k+1}^*(f_k(x_k, u_k, w_k)) \right\}, \quad \text{for all } x_k.$$

- The optimal cost $J^*(x_0)$ is obtained at the last step: $J_0^*(x_0) = J^*(x_0)$.

Online implementation of the optimal policy, given J_1^*, \dots, J_{N-1}^*

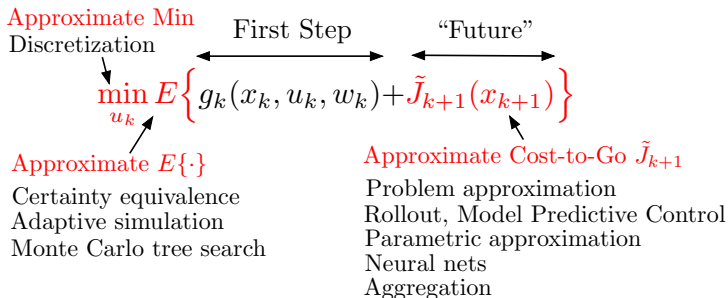
Sequentially, going forward, for $k = 0, 1, \dots, N - 1$, observe x_k and apply

$$u_k^* \in \arg \min_{u_k \in U_k(x_k)} E \left\{ g_k(x_k, u_k, w_k) + J_{k+1}^*(f_k(x_k, u_k, w_k)) \right\}.$$

The main difficulties: Too much computation, too much memory storage.

Approximation in value space: Use \tilde{J}_k in place of J_k^* ; possibly approximate $E\{\cdot\}$ and \min_{u_k} .

Approximation in Value Space: One-Step Lookahead




Approximation in value space uses \tilde{J}_{k+1} (in place of J_{k+1}^*) and lookahead minimization, to construct suboptimal control law $\tilde{\mu}_k$ at time k .

Three main Issues; they can be addressed separately

- How to construct \tilde{J}_k , $k = 1, \dots, N$.
- How to simplify $E\{\cdot\}$ operation.
- How to simplify min operation.

Approximation in Value Space: Multistep Lookahead



- At state x_k , we solve an ℓ -stage version of the DP problem with x_k as the initial state and $\tilde{J}_{k+\ell}$ as the terminal cost function.
- Use the first control of the ℓ -stage policy thus obtained, while discarding the others.

Can view ℓ -step lookahead as a special case of one-step lookahead:

The “effective” one-step lookahead function is the optimal cost function of an $(\ell - 1)$ -stage DP problem with terminal cost $\tilde{J}_{k+\ell}$.

Approximation in Policy Space: The Major Alternative to Approximation in Value Space

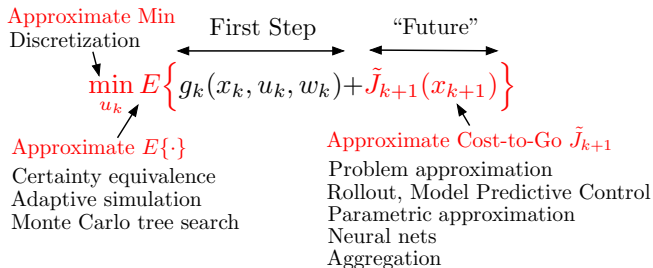
- Idea: Select the policy by **optimization over a suitably restricted class of policies**.
- The restricted class is usually a parametric family of policies $\mu_k(x_k, r_k)$, $k = 0, \dots, N - 1$, of some form, where r_k is a parameter (e.g., a neural net).
- **Important advantage once the parameters r_k are computed**: The computation of controls during on-line operation of the system is often much easier: At state x_k apply $u_k = \mu_k(x_k, r_k)$.

Approximation in policy space on top of approximation in value space

- Compute approximate cost-to-go functions \tilde{J}_{k+1} , $k = 0, \dots, N - 1$.
- This defines the corresponding suboptimal policy $\tilde{\mu}_k$, $k = 0, \dots, N - 1$, through one-step or multistep lookahead.
- **Approximate $\tilde{\mu}_k$ using some form of regression** and a training set consisting of a large number q of sample pairs (x_k^s, u_k^s) , $s = 1, \dots, q$, where $u_k^s = \tilde{\mu}_k(x_k^s)$.
- **Example**: Introduce a parametric family of policies $\mu_k(x_k, r_k)$, $k = 0, \dots, N - 1$, of some form, where r_k is a parameter. Then estimate the parameters r_k by

$$r_k \in \arg \min_r \sum_{s=1}^q \|u_k^s - \mu_k(x_k^s, r)\|^2.$$

On-Line and Off-Line Lookahead Implementations



- For many-state problems, **the minimizing controls $\tilde{\mu}_k(x_k)$ are computed on-line** (storage issue).
- **Off-line methods:** All the functions \tilde{J}_{k+1} are computed for every k , before the control process begins.
- **Examples of off-line methods:** Neural network and other parametric approximations; also aggregation.
- **On-line methods:** The values $\tilde{J}_{k+1}(x_{k+1})$ are computed only at the relevant next states x_{k+1} , and are used to compute the control to be applied at the N time steps.
- **Examples of on-line methods:** Rollout and model predictive control.
- **Rollout is well-suited for on-line replanning**, but lots of on-line computation.

Simplifying the Minimization of the Expected Value in Lookahead Schemes

$$\min_{u_k \in U_k(x_k)} E\left\{g_k(x_k, u_k, w_k) + \tilde{J}_{k+1}(f_k(x_k, u_k, w_k))\right\}$$

- If $U_k(x_k)$ is a finite set, the minimization can be done by **brute force**.
- If $U_k(x_k)$ is an infinite set, it may be replaced by a finite set through **discretization**.
- For deterministic problems and continuous control spaces, a more efficient alternative may be to use **nonlinear programming** techniques.
- For stochastic problems and continuous control spaces, we may use **stochastic programming**. Lookahead must be short because of the high branching factor of the lookahead tree when the problem is stochastic.

One possibility to deal with the $E\{\cdot\}$:

Assumed certainty equivalence, i.e., choose a typical value \tilde{w}_k of w_k , and use the control $\tilde{u}_k(x_k)$ that solves the deterministic problem

$$\min_{u_k \in U_k(x_k)} \left[g_k(x_k, u_k, \tilde{w}_k) + \tilde{J}_{k+1}(f_k(x_k, u_k, \tilde{w}_k)) \right]$$

However, this may degrade performance significantly.

Model-Based Versus Model-Free Implementation

Our layman's use of the term “**model-free**”: A method is called model-free if it involves calculations of expected values using **Monte Carlo simulation**.

Model-free is necessary when:


- A **mathematical model of the probabilities $p_k(w_k | x_k, u_k)$ is not available** but a computer model/simulator is. For any (x_k, u_k) , it simulates sample probabilistic transitions to a successor state x_{k+1} , and generates the corresponding transition costs.
- When for reasons of computational efficiency **we prefer to compute the expected value by using sampling** and Monte Carlo simulation; e.g., approximate an integral or a huge sum of numbers by a Monte Carlo estimate.

Principal example: Calculations of approximate Q-factors in lookahead schemes

$$E\left\{g_k(x_k, u_k, w_k) + \tilde{J}_{k+1}(f_k(x_k, u_k, w_k))\right\}$$

(**assuming \tilde{J}_{k+1} has been computed**).

Model-Free Q-Factor Calculation for Stochastic Problems



- Use the simulator to collect a large number of “representative” samples of state-control-successor states-stage cost quadruplets $(x_k^s, u_k^s, x_{k+1}^s, g_k^s)$, and corresponding sample Q-factors

$$\beta_k^s = g_k^s + \tilde{J}_{k+1}(x_{k+1}^s), \quad s = 1, \dots, q$$

- Introduce a parametric family of Q-factors $\tilde{Q}_k(x_k, u_k, r_k)$.
- Determine the parameter vector \bar{r}_k by the least-squares regression

$$\bar{r}_k \in \arg \min_{r_k} \sum_{s=1}^q (\tilde{Q}_k(x_k^s, u_k^s, r_k) - \beta_k^s)^2$$

- Use the policy

$$\tilde{\mu}_k(x_k) \in \arg \min_{u_k \in U_k(x_k)} \tilde{Q}_k(x_k, u_k, \bar{r}_k)$$

Multistep Lookahead Issues

DP minimization

↓

$$\min_{u_k, \mu_{k+1}, \dots, \mu_{k+\ell-1}} E \left\{ g_k(x_k, u_k, w_k) + \sum_{m=k+1}^{k+\ell-1} g_m(x_m, \mu_m(x_m), w_m) + \tilde{J}_{k+\ell}(x_{k+\ell}) \right\}$$

First ℓ Steps

“Future”

Lookahead Minimization

Cost-to-go Approximation


Main hope

- Minimization over many steps will work better than minimization over few steps (with long enough lookahead we are optimal).
- By using a long-step lookahead, we can afford a simpler/less accurate cost-to-go approximation.

Main Issue

Minimization over many stages is costly; stochastic problems are harder because of a larger branching factor of the lookahead tree.

Multistep Lookahead and Deterministic Problems



If the problem is **deterministic and finite-state**, the lookahead minimization is a **shortest path problem** and may be solved on-line.

If the problem is **deterministic and continuous-state/control**, the lookahead minimization may be quickly solvable by **nonlinear programming** (model predictive control case).

If the problem is **stochastic and finite-state**, the lookahead minimization can be **split into a first stochastic step and a deterministic remainder**; i.e., use a deterministic shortest path problem approximation for the remaining steps.

Let's Take a Working Break to Consider the Following Challenge Questions

Question 1

Consider one-step lookahead with two different cost function approximations \tilde{J}_{k+1} and \hat{J}_{k+1} . Assume that \tilde{J}_{k+1} is “much closer” to J_{k+1}^* (in any way you may think of).


Will \tilde{J}_{k+1} produce a better policy than \hat{J}_{k+1} ?

Question 2


Consider multistep lookahead.

Will longer lookahead produce a better policy than shorter lookahead?

The Answers are NO and NO



Constant shift in Q-factor does not affect the minimizing control. For a good suboptimal policy, the “slope” of the Q-factor difference $Q_k(x_k, u) - \tilde{Q}_k(x_k, u)$ should be small.



Problem with “edge effects”: Two controls, u (optimal) and u' (suboptimal), and cost function approximation $\tilde{J}_k(x_k) \equiv 0$. u will be preferred based on ℓ -step lookahead. u' will be preferred based on $(\ell + 1)$ -step lookahead.


$$\tilde{\mu}_k(x_k) \in \arg \min_{u_k \in U_k(x_k)} E \left\{ g_k(x_k, u_k, w_k) + \tilde{J}_{k+1}(f_k(x_k, u_k, w_k)) \right\}$$

\tilde{J}_{k+1} is the optimal cost function of a simpler problem
How this is done is typically problem-dependent

Examples


- **Enforced decomposition** of systems that consist of weakly coupled subsystems. Leads to single subsystem computations.
- **Probabilistic approximation**. Enforced certainty equivalence. Leads to deterministic minimizations with no expected values to compute.
- **Aggregation**. Construct a “smaller” aggregate problem by introducing aggregate states.

Example: Optimize the Routes of n Vehicles Through a Road Network



- Aim: **Execute a number of tasks with given values**
- The value of a task is collected only once; a finite horizon is assumed.
- This is a very complex combinatorial problem.
- The single vehicle problem is typically much simpler (e.g., can be solved exactly or with a high-quality heuristic).
- At a given state: **Solve (suboptimally) the tail subproblem one-vehicle-at-a-time.**

Enforced Decomposition: Constraint Decoupling by Relaxation




- Let $\mathbf{x}_k = (x_k^1, \dots, x_k^n)$, $\mathbf{u}_k = (u_k^1, \dots, u_k^n)$, $\mathbf{w}_k = (w_k^1, \dots, w_k^n)$, with (x_k^i, u_k^i, w_k^i) corresponding to the i th subsystem.
- Assume that the only coupling between subsystems is the control constraint

$$(u_k^1, \dots, u_k^n) \in U, \quad \text{e.g., } u_k^i \in U^i, \quad u_k^1 + \dots + u_k^n \leq b_k$$

- **Approximate U with a decomposed constraint $U^1 \times \dots \times U^n$.**
- **The problem decomposes into n decoupled subproblems.** Let \tilde{J}_k^i be the optimal cost to go functions for the i th decoupled subproblem (obtained by DP off-line).
- Use one-step lookahead with the cost-to-go approximation

$$\tilde{J}_{k+1}(\mathbf{x}_{k+1}) = \tilde{J}_{k+1}^1(x_{k+1}^1) + \dots + \tilde{J}_{k+1}^n(x_{k+1}^n)$$

Example: Production Planning



A work center producing n product types

- x_k^i, u_k^i, w_k^i : the amounts stored, produced, and demanded of product i at time k
- State is the stock vector $x_k = (x_k^1, \dots, x_k^n)$, where $x_{k+1}^i = x_k^i + u_k^i - w_k^i$
- U represents the (shared) production capacity of the work center
- In a more complex version (involving equipment failures), U depends on a random parameter α_k that changes according to a Markov chain

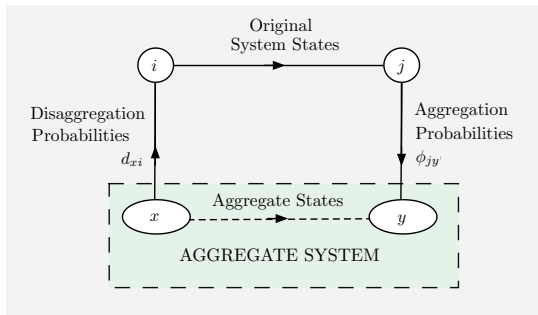
Another example in the text: “**Restless**” multiarmed bandit problems. (Select one out of n projects to work on at each stage.)

Modify the probability distributions $P(w_k \mid x_k, w_k)$ to simplify the calculation of $\tilde{J}_{k+\ell}$ and/or the lookahead minimization.

Certainty equivalent control (inspired by linear-quadratic control problems)

- Replace uncertain quantities with deterministic nominal values.
- The lookahead and tail problems are deterministic so they could be solvable by DP or by special deterministic methods on-line.
- Use expected values or forecasts to determine nominal values; update policy when forecasts change (on-line replanning).
- Use state estimates instead of belief states.
- A variant: Partial certainty equivalence. Fix only some uncertain quantities to nominal values.
- A generalization: Approximate $E\{\cdot\}$ by limited simulation.

Problem Approximation by Aggregation (to be discussed in detail later)



- Construct a “smaller” aggregate problem by introducing aggregate states.
- Use the exact costs-to-go of the aggregate tail problem as approximate costs-to-go for the original.

Aggregation examples:

- State discretization-interpolation schemes.
- Grouping of states into subsets, which serve as aggregate states.
- Feature-based discretization; aggregate problem operates in the space of features.

We will cover:

- Rollout for deterministic and stochastic problems
- Monte Carlo tree search
- Model predictive control

PLEASE READ AS MUCH OF SECTIONS 2.3, 2.4 AS YOU CAN
PLEASE DOWNLOAD THE LATEST VERSIONS FROM MY WEBSITE