

Reinforcement Learning and Optimal Control


ASU, CSE 691, Winter 2019

Dimitri P. Bertsekas
dimitrib@mit.edu

Lecture 6

- 1 Parametric Approximation Architectures
- 2 Training of Approximation Architectures
- 3 Incremental Optimization of Sums of Differentiable Functions
- 4 Neural Networks
- 5 Neural Nets and Finite Horizon DP

Recall the Approximation in Value Space Framework for Finite Horizon Problems



An **approximation architecture** is a class of functions $\tilde{J}(x, r)$ that depend on x and a vector $r = (r_1, \dots, r_m)$ of m “tunable” scalar parameters (or weights).

Issues and terminology

- Aim: Choose r to **make $\tilde{J}(x, r)$ close to some target cost function $J(x)$** .
- Training algorithm chooses r . It typically uses **least squares optimization** (regression) to fit $\tilde{J}(x, r)$ to a data set of state-cost pairs.
- An architecture is called **linear** if $\tilde{J}(x, r)$ is linear in r .
- It is called **feature-based** if it depend on x via a feature vector $\phi(x)$,

$$\tilde{J}(x, r) = \hat{J}(\phi(x), r),$$

where \hat{J} is some function. Idea: **Features capture dominant nonlinearities**.

- A **linear feature-based architecture**:

$$\tilde{J}(x, r) = \sum_{\ell=1}^m r_{\ell} \phi_{\ell}(x) = r' \phi(x),$$

where r_{ℓ} and $\phi_{\ell}(x)$ are the ℓ th components of r and $\phi(x)$.

Least squares regression

- Collect a set of state-cost training pairs (x^s, β^s) , $s = 1, \dots, q$, where β^s is equal to the target cost $J(x^s)$ plus some “noise”.
- r is determined by solving the problem

$$\min_r \sum_{s=1}^q (\tilde{J}(x^s, r) - \beta^s)^2$$

- Sometimes a quadratic regularization term $\gamma \|r\|^2$ is added to the least squares objective, to facilitate the minimization (among other reasons).

Training of linear feature-based architectures can be done exactly

- If $\tilde{J}(x, r) = r' \phi(x)$, where $\phi(x)$ is the m -dimensional feature vector, the training problem is quadratic and can be solved in closed form.
- The exact solution of the training problem is given by

$$\hat{r} = \left(\sum_{s=1}^q \phi(x^s) \phi(x^s)' \right)^{-1} \sum_{s=1}^q \phi(x^s) \beta^s$$

- This requires a lot of computation for a large m and data set; may not be best.

The main training issue

How to exploit the structure of the training problem

$$\min_r \sum_{s=1}^q (\tilde{J}(x^s, r) - \beta^s)^2$$

to solve it efficiently.

Key characteristics of the training problem

- **Possibly nonconvex with many local minima**, horribly complicated graph of the cost function (true when a neural net is used).
- **Many terms in the least squares sum**; standard gradient and Newton-like methods are essentially inapplicable.
- **Incremental** iterative methods that operate on **a single term** $(\tilde{J}(x^s, r) - \beta^s)^2$ **at each iteration** have worked well enough (for many problems).

Generic sum of terms optimization problem

Minimize

$$f(y) = \sum_{i=1}^m f_i(y)$$

where each f_i is a differentiable scalar function of the n -dimensional vector y (this is the parameter vector in the context of parametric training).

The ordinary gradient method generates y^{k+1} from y^k according to

$$y^{k+1} = y^k - \gamma^k \nabla f(y^k) = y^k - \gamma^k \sum_{i=1}^m \nabla f_i(y^k)$$

where $\gamma^k > 0$ is a stepsize parameter.


The incremental gradient counterpart

Choose an index i_k and iterate according to

$$y^{k+1} = y^k - \gamma^k \nabla f_{i_k}(y^k)$$

where $\gamma^k > 0$ is a stepsize parameter.

The Advantage of Incrementalism



$$\text{Minimize } f(y) = \frac{1}{2} \sum_{i=1}^m (c_i y - b_i)^2$$

Compare the ordinary and the incremental gradient methods in two cases

- When far from convergence: **Incremental gradient is as fast as ordinary gradient with $1/m$ amount of work.**
- When close to convergence: **Incremental gradient gets confused** and requires a diminishing stepsize for convergence.

Incremental **aggregated** method aims at acceleration

- Evaluates gradient of a single term at each iteration.
- Uses previously calculated gradients as if they were up to date

$$y^{k+1} = y^k - \gamma^k \sum_{\ell=0}^{m-1} \nabla f_{i_{k-\ell}}(y^{k-\ell})$$

- Has theoretical and empirical support, and it is often preferable.

Stochastic gradient method (also called stochastic gradient descent or **SGD**)

- Applies to **minimization** of $f(y) = E\{F(y, w)\}$ where w is a random variable
- Has the form

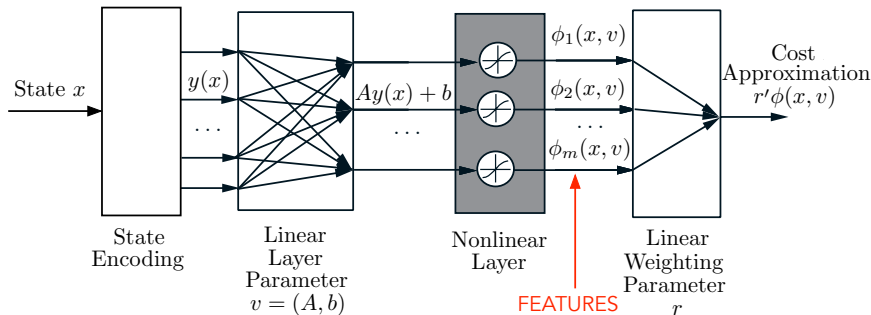
$$y^{k+1} = y^k - \gamma^k \nabla_y F(y^k, w^k)$$

where w^k is a sample of w and $\nabla_y F$ denotes gradient of F with respect to y .

- **The incremental gradient method with random index selection is the same as SGD** [convert the sum $\sum_{i=1}^m f_i(y)$ to an expected value, where i is random with uniform distribution].

- How to pick the **stepsize** γ^k (usually $\gamma^k = \frac{\gamma}{k+1}$ or similar).
- How to deal (if at all) with **region of confusion** issues (detect being in the region of confusion and reduce the stepsize).
- How to select the **order of terms to iterate** (cyclic, random, other).
- **Diagonal scaling** (a different stepsize for each component of y).
- **Alternative methods** (more ambitious): Incremental Newton method, extended Kalman filter (see the textbook and references).

Neural Nets: An Architecture that Automatically Constructs Features




Given a set of state-cost training pairs (x^s, β^s) , $s = 1, \dots, q$, the parameters of the neural network (A, b, r) are obtained by solving the training problem


$$\min_{A, b, r} \sum_{s=1}^q \left(\sum_{\ell=1}^m r_{\ell} \sigma \left((Ay(x^s) + b)_{\ell} \right) - \beta^s \right)^2$$

- Incremental gradient is typically used for training.
- **Universal approximation property.**

Rectifier and Sigmoidal Nonlinearities




The **rectified linear unit** $\sigma(\xi) = \ln(1 + e^\xi)$. It is the rectifier function $\max\{0, \xi\}$ with its corner “smoothed out.”




Sigmoidal units: The **hyperbolic tangent** function $\sigma(\xi) = \tanh(\xi) = \frac{e^\xi - e^{-\xi}}{e^\xi + e^{-\xi}}$ is on the left. The **logistic** function $\sigma(\xi) = \frac{1}{1 + e^{-\xi}}$ is on the right.

Deep Neural Networks




- The multilayer network provides **a hierarchy of features** (each set of features being a function of the preceding set of features).
- We may **use matrices A with a special structure** that encodes special linear operations such as convolution.
- When such structures are used, the training problem may become easier, because the number of parameters in the linear layers is drastically decreased.
- They have been found **more effective than shallow neural nets** for some problems.
- **Incremental gradient is still used for training.** The algorithm is based on **an intelligent way of using the chain rule** to calculate the incremental gradient at each iteration.

A Working Break: Challenge Question





How can we use linear and rectifier units to **construct the “pulse” feature below?**



- What are the features that can be produced by neural nets?
- Why do neural nets have a “universal approximation” property?

Answer



Using the pulse feature as a building block, any feature can be approximated

Sequential DP Approximation - A Parametric Approximation at Every Stage (Also Called **Fitted Value Iteration**)

Start with $\tilde{J}_N = g_N$ and **sequentially train going backwards**, until $k = 0$

- Given a cost-to-go approximation \tilde{J}_{k+1} , we **use one-step lookahead to construct a large number of state-cost pairs** (x_k^s, β_k^s) , $s = 1, \dots, q$, where

$$\beta_k^s = \min_{u \in U_k(x_k^s)} E \left\{ g(x_k^s, u, w_k) + \tilde{J}_{k+1}(f_k(x_k^s, u, w_k), r_{k+1}) \right\}, \quad s = 1, \dots, q$$

- We “train” an architecture \tilde{J}_k on the training set (x_k^s, β_k^s) , $s = 1, \dots, q$.

Typical approach: Train by least squares/regression and possibly using a neural net

We minimize over r_k

$$\sum_{s=1}^q (\tilde{J}_k(x_k^s, r_k) - \beta_k^s)^2$$

Sequential Q-Factor Approximation

- Consider sequential DP approximation of Q -factor parametric approximations

$$\tilde{Q}_k(x_k, u_k, r_k) = E \left\{ g_k(x_k, u_k, w_k) + \min_{u \in U_{k+1}(x_{k+1})} \tilde{Q}_{k+1}(x_{k+1}, u, r_{k+1}) \right\}$$

(Note a mathematical magic: **The order of $E\{\cdot\}$ and min have been reversed.**)

- We obtain $\tilde{Q}_k(x_k, u_k, r_k)$ by training with many pairs $((x_k^s, u_k^s), \beta_k^s)$, where **β_k^s is a sample of the approximate Q -factor of (x_k^s, u_k^s) .** [No need to compute $E\{\cdot\}$.]
- Note: **No need for a model to obtain β_k^s .** Sufficient to have a simulator that generates state-control-cost-next state random samples

$$((x_k, u_k), (g_k(x_k, u_k, w_k), x_{k+1}))$$

- Having computed r_k , the one-step lookahead control is obtained on-line as

$$\bar{\mu}_k(x_k) \in \arg \min_{u \in U_k(x_k)} \tilde{Q}_k(x_k, u, r_k)$$

without the need of a model or expected value calculations.

- Important advantage: The **on-line calculation of the control is simplified.**

We will cover:

- Infinite horizon DP problems: Stochastic shortest path and discounted problems
- Analysis, Bellman's equation, optimality conditions
- Algorithms: Value iteration, policy iteration
- We will likely need more than one lecture

PLEASE READ AS MUCH OF SECTIONS 4.1-4.5 AS YOU CAN
APPENDIX OF CHAPTER 4 CONTAINS PROOFS; TAKE A CRACK AT THEM
PLEASE DOWNLOAD THE LATEST VERSIONS FROM MY WEBSITE